COMP550 Final Project Writeup

Anonymous ACL submission

Abstract

We tried to replicate results on question answering data from big biomedical models like BioGPT or BioBERT while using far less training data. We fine-tuned three models, GPT-2, BERT, and T5, on QA data and used retrieval augmented generation techniques to improve their performance.

1 Introduction

Large Language Models (LLMs) benefit from the fact that the number of parameters they have is well above millions. Some even reach billions of parameters. They also benefit from the amount of data they are able to use during pre-training. This makes the models competent at a wide range of tasks.

Typical LLMs, used by many globally, are extremely well established at question answering (QA). The amount of knowledge they have on a number of different subjects makes them capable of answering most questions, save for some extremely difficult PhD level ones.

The biomedical field is one of the more difficult ones for the average LLM because of the amount of data in the field as well as the interactions between different parts of it. For example, drugs interact with diseases which have symptoms which may be side effects of other drugs or other diseases altogether. The model must be able to, given a question, filter through this information and deliver a clear and concise answer. In this project, we use datasets sampled from medical QA tests like the USMLE (Jin et al., 2020), a dataset pulled directly from medical exams. We also use MedMCQA (Pal et al., 2022), another multiple-choice QA dataset pulled from exams. These datasets include the types of questions that would stump smaller or less expressive models.

The USMLE dataset includes questions about choosing the correct condition given a number of symptoms and a patient's information. MedMCQA contains more textbook information, including facts about treatments, drugs, pathologies, and more.

Larger LLMs are able to answer these questions well above chance level (Liévin et al., 2023), however smaller models, like BERT (Devlin et al., 2019), are still in the 30% range on these 4 option questions.

We explore whether models trained explicitly on biomedical question-answer data can come close to outperforming large models that are trained on a massive amount of general data. We attempt to use techniques in order to improve training as well as inference with as little cost as possible.

2 Related Work

Numerous models have already attempted to train on large biomedical corpora such as PubMed. One such model, PubMedBERT is based on BERT and pretrained from scratch on unlabeled biomedical data, mostly PubMed papers and abstracts (Gu et al., 2021). Other models have been focused on training on biomedical data as a whole, as opposed to just papers. These models also perform well on the datasets described in the previous section. BioGPT (Luo et al., 2022) was a response to BioBERT (Lee et al., 2019) which used GPT's generative abilities to improve performance.

All of these models are pre-trained from scratch on millions (or more) of passages from biomedical textbooks, papers, articles, and more. Thus, they are trained on multiple GPUs for longer than we could afford to wait. In order to improve training times, as well as accuracy, we decided to implement retrieval augmented generation (RAG) (Lewis et al., 2021), a technique in which the model retrieves data based on the prompt which it can use as context for the given task. It is a related but different idea to in-context learning. With RAG, the models are injecting explicit data that they can use for inference, while with ICL, this data is implicit to the prompt.

The benefit of RAG is that the model need not train on all the data but can instead just pull valuable passages that it can use to come up with an answer. This has previously been done in the biomedical field, starting with MedRAG (Xiong et al., 2024), an RAG framework that has been shown to improve the accuracy of GPT3.5 on most tasks, including the QA we look at. It is impor-

tant to note that both the dataset the contexts are pulled from, as well as the retriever, are important in improving the model's performance. For example, Wikipedia's general knowledge may not help in USMLE style questions where knowing how conditions manifest matters. In that case, case studies seen in PubMed may be better.

Further, research has shown that fine-tuning RAG-augmented models may improve their accuracy even further (Hassan et al., 2025).

3 Methods

3.1 GPT

3.2 BERT

BERT (Devlin et al., 2019) is an encoder-only transformer. This project aims to compare this general-purpose model and BioBERT (Lee et al., 2019) which was pre-trained on biomedical data text.

For multiple-choice QA, each (question, answer) pair is encoded independently as `[CLS]` q `[SEP]` a_i `[SEP]`. The classification token embedding is then passed through a linear head to produce a score per option, which are are normalised with softmax and trained via cross-entropy. To avoid the cost of full finetuning, we used LoRA (Hu et al., 2021) with a rank of 32, which drastically reduces trainable parameters while maintaining comparable performance. Hyperparameters (learning rate $= 3\times10^{-4}$ and batch size $= 8$) were selected via a one-epoch grid search using the MedMCQA validation accuracy. Using the combination with the highest validation accuracy, models were finetuned on a shuffled combination of MedMCQA and USMLE, all with four options.

In order to increase accuracy, we also trained RAG (Lewis et al., 2021) variants, so the model learns from both the QA patterns and from the retrived medical knowledge. For each training example, the top three passages retrieved by the BM25 index (Lewis et al., 2021) were prepended to the question before tokenisation, producing inputs of the form: `[CLS]` c_1 c_2 c_3 q `[SEP]` a_i `[SEP]` with a 512-token limit. The retrieval corpus consisted of two MedRAG corpora: the medical textbooks corpora and a 150k sample of the Wikipedia corpus. Because the raw Wikipedia corpus contains a large proportion of non-medical content, it was first filtered with common medical terms. All passages were then chunked into 100-word segments. At query time, the question text is tokenised and the top-k passages with the highest BM25 scores are returned. We used $k = 3$.

The final model was tested on MedMCQA, USMLE and GPQA.

3.3 T5

T5 is an encoder-decoder developed by Google. Its intruction-tuned version, Flan-T5 (Chung et al., 2022) is what was used for this part of the project. The version used was flan-t5-xl, the largest available model. The reason for using this model was twofold: it would be most probable to get the best baseline accuracy due to its size, and it would be less likely to overfit (again due to its size). In this case, overfitting meant collapsing to selecting the same answer on multiple choice questions.

Further, this part of the project used a FAISS (Douze et al., 2025) index of MedRAG's Wikipedia, textbooks, and StatPearl corpora. The textbooks and StatPearl corpora were quite small, so they were used in full, while only 150k Wikipedia samples were chosen at random. These were then chunked into 100-word passages and indexed using FAISS. All in all, this came out to only 2GB of data; far less than the entirety of MedRAG, which is over 100GB.

To save time and memory, the model was allowed to use the top 3 contexts scored by the MedCPT retriever (Jin et al., 2023). Then, the model was trained on QA data with RAG. That is, the model was allowed to train by answering retrieval-augmented questions, as opposed to just training on the questions alone. The idea was to allow the model to learn patterns in the questions, as well as learn information from the contexts retrieved from the FAISS index.

Because T5-xl is an extra large model, we furthered our efforts to save memory. Thus, we used 4-bit quantization (Dettmers et al., 2023) and Low-Rank Adaptation (LoRA) (Hu et al., 2021). The quantization technique decreases the precision of the floating points used from default (16-bit in this case) to 4-bit. This saves memory at the cost of some accuracy, but it has been shown that accuracy does not suffer as much as expected. The last of our efforts used LoRA, which only uses a small subset of trainable parameters. The idea is that the updated weights are typically of low rank, so instead of training the full model and only updating low-rank weights, we only train with the low-rank weights. This makes the model much easier to fit in memory and decreases training time substantially.

This fine-tuned model was then tested on MedMCQA, USMLE, and GPQA.

4 Results

4.1 GPT

4.2 BERT

Table 1 reports accuracy across all three benchmarks. On both MedMCQA and USMLE, fine-tuning produced clear gains over the random base-

line of 25%, and in both cases RAG-augmented training improved over standard fine-tuning, suggesting that prepending retrieved passages during training provides useful signal beyond the QA data alone. As for BioBERT without any fine-tuning, it consistently scored below the random baseline and below fine-tuned BERT, suggesting that biomedical pre-training alone, without task-specific fine-tuning, is insufficient for multiple-choice QA. On GPQA-Diamond, all three models fell below the random baseline.

Table 1: Accuracy on MedMCQA, USMLE and GPQA Random chance is 25% for all benchmarks.

Model	MedMCQA	USMLE	GPQA-Diamond
BERT (standard)	32.11%	27.34%	22.22%
BERT (RAG)	36.10%	28.99%	24.75%
BioBERT	24.96%	24.74%	22.22%
Random baseline	25.00%	25.00%	25.00%

4.3 T5

The baseline accuracies for the T5 model are shown in table 2. Immediately, it is noticeable that USMLE is a harder dataset. This is due to the fact that instead of exam-like questions, they are more case studies. As mentioned in the introduction, these questions ask the student to correctly identify a condition based on some symptoms and patient factors.

Method	MedMCQA	USMLE
Pre-FT Baseline	0.282	0.280
Pre-FT RAG	0.322	0.273

Table 2: T5 accuracy on test datasets before fine-tuning.

The model was then fine-tuned on balanced split over all labels from both USMLE and MedMCQA. In total, the model was trained on 22696 examples, including the question *and* contexts. See figure 1 for the loss, computed every 300 steps. To make sure the model was not collapsing, we tracked the number of outputs of each label. These did not show any crazy results and will not be shown.

Finally, the model was retested on the exact same test set as before fine-tuning. The results are shown in table 3.

Method	MedMCQA	USMLE
Post-FT Baseline	0.236	0.225
Post-FT RAG	0.239	0.227

Table 3: T5 accuracy on test datasets after fine-tuning.

As a last resort, we checked whether the prompt was an issue. Thus, we tried including the options into the FAISS context retrieval and achieved the results in table 4.

Dataset	Accuracy
USMLE (Post-FT RAG, q+options)	0.228
MedMCQA (Post-FT RAG, q+options)	0.290

Table 4: Test set accuracy for Post-FT + RAG (question + options)

5 Discussion & Conclusion

BERT (standard) improved over the random baseline by roughly 7% on MedMCQA and 2% on USMLE, and RAG-augmented training added 4% and 2% respectively. This suggests that retrieved context provides complementary knowledge to that learned during finetuning. However, the results also expose the ceiling of encoder-only models for this task. GPQA scores below random chance for all three models. This proves that GPQA questions require advanced deductive reasoning, which cannot be performed by a model that compresses its entire input into a single `[CLS]` vector and then scores options. Prepending retrieved passages helps when the answer can be read directly from the context, but is ineffective when reasoning is required. As for the BioBERT results, despite being pre-trained on biomedical text, it BioBERT performs below the random baseline on all three benchmarks. This indicates that domain-specific pre-training does not transfer to multiple-choice QA without task-specific fine-tuning. Overall, these findings suggest that encoder-only models are not well suited to graduate-level biomedical reasoning regardless of fine-tuning or retrieval augmentation.

Our T5 model had an accuracy increase of 4% on the MedMCQA test set but actually dropped in accuracy on the USMLE dataset when only using RAG (no fine-tuning). This is likely due to the more reasoning-heavy nature of USMLE. However, after fine-tuning, the model performed worse on both datasets overall, dipping below chance. Originally, this was caused by a collapse in the model's selection of labels (e.g: always selecting C). After fixing this, the model still performs just as poorly.

The drop in performance could be due to overfitting, however the train and validation loss do not corroborate this. It would have been better to track accuracy on the training and validation set, but this would've increased training time too much.

After looking at the poor retrievals for certain questions, we opted to use the same evaluation, except for including the options into the FAISS retrieval. These results show a massive improvement on the fine-tuned model on MedMCQA and a small increase on USMLE for the same model. In

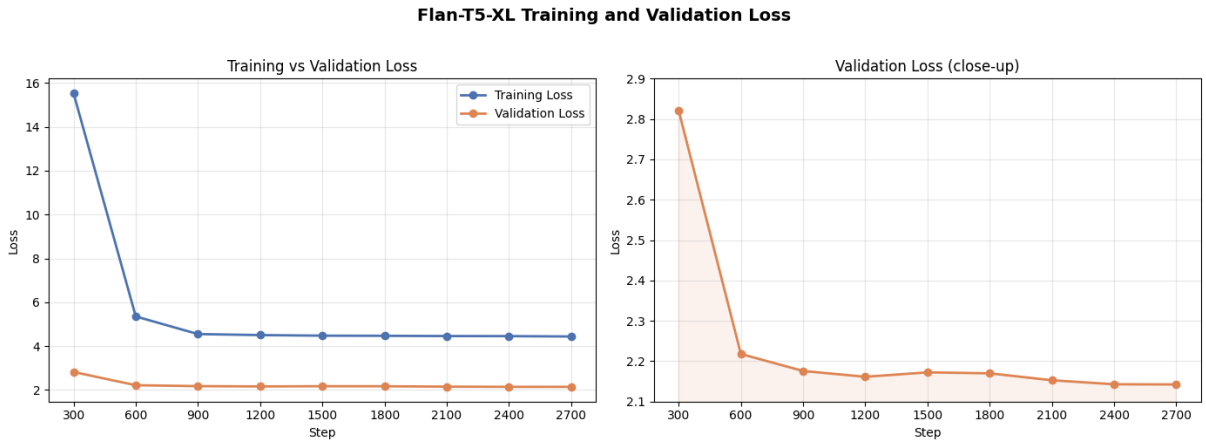

Figure 1: Training and validation loss visualization of flan-T5-xl

all likelihood, the contexts were insufficient during training, as some questions are of the "which of the below are correct" sort. Thus, including the options would help. In the future, using better prompting and FAISS retrieval would be a focus.

6 Statement of Contributions

All preprocessing, training, and testing was done separately. James worked with GPT, Jessica worked with BERT, and Tyler worked with T5.

References

Hyung Won Chung, Le Hou, Shayne Longpre, Barret Zoph, Yi Tay, William Fedus, Yunxuan Li, Xuezhi Wang, Mostafa Dehghani, Siddhartha Brahma, Albert Webson, Shixiang Shane Gu, Zhuyun Dai, Mirac Suzgun, Xinyun Chen, Aakanksha Chowdhery, Alex Castro-Ros, Marie Pellat, Kevin Robinson, and 16 others. 2022. Scaling instruction-finetuned language models. *Preprint*, arXiv:2210.11416.

Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. 2023. Qlora: Efficient finetuning of quantized llms. *Preprint*, arXiv:2305.14314.

Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. Bert: Pre-training of deep bidirectional transformers for language understanding. *Preprint*, arXiv:1810.04805.

Matthijs Douze, Alexandr Guzhva, Chengqi Deng, Jeff Johnson, Gergely Szilvasy, Pierre-Emmanuel Mazaré, Maria Lomeli, Lucas Hosseini, and Hervé Jégou. 2025. The faiss library. *Preprint*, arXiv:2401.08281.

Yu Gu, Robert Tinn, Hao Cheng, Michael Lucas, Naoto Usuyama, Xiaodong Liu, Tristan Naumann, Jianfeng Gao, and Hoifung Poon. 2021. Domain-specific language model pretraining for biomedical natural language processing. *ACM Transactions on Computing for Healthcare*, 3(1):1–23.

Tasnimul Hassan, Md Faisal Karim, Haziq Jeelani, Elham Behnam, Robert Green, and Fayeq Jeelani Syed. 2025. Optimizing medical question-answering systems: A comparative study of fine-tuned and zero-shot large language models with rag framework. *Preprint*, arXiv:2512.05863.

Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2021. Lora: Low-rank adaptation of large language models. *Preprint*, arXiv:2106.09685.

Di Jin, Eileen Pan, Nassim Oufattole, Wei-Hung Weng, Hanyi Fang, and Peter Szolovits. 2020. What disease does this patient have? a large-scale open domain question answering dataset from medical exams. *arXiv preprint arXiv:2009.13081*.

Qiao Jin, Won Kim, Qingyu Chen, Donald C Comeau, Lana Yeganova, W John Wilbur, and Zhiyong Lu. 2023. Medcpt: Contrastive pre-trained transformers with large-scale pubmed search logs for zero-shot biomedical information retrieval. *Bioinformatics*, 39(11).

Jinhyuk Lee, Wonjin Yoon, Sungdong Kim, Donghyeon Kim, Sunkyu Kim, Chan Ho So, and Jaewoo Kang. 2019. Biobert: a pre-trained biomedical language representation model for biomedical text mining. *Bioinformatics*, 36(4):1234–1240.

Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen tau Yih, Tim Rocktäschel, Sebastian Riedel, and Douwe Kiela. 2021. Retrieval-augmented generation for knowledge-intensive nlp tasks. *Preprint*, arXiv:2005.11401.

Valentin Liévin, Christoffer Egeberg Hother, Andreas Geert Motzfeldt, and Ole Winther. 2023.

Can large language models reason about medical questions? *Preprint*, arXiv:2207.08143.

Renqian Luo, Liai Sun, Yingce Xia, Tao Qin, Sheng Zhang, Hoifung Poon, and Tie-Yan Liu. 2022. Biogpt: generative pre-trained transformer for biomedical text generation and mining. *Briefings in Bioinformatics*, 23(6).

Ankit Pal, Logesh Kumar Umapathi, and Malaikannan Sankarasubbu. 2022. Medmcqa: A large-scale multi-subject multi-choice dataset for medical domain question answering. In *Proceedings of the Conference on Health, Inference, and Learning*, volume 174 of *Proceedings of Machine Learning Research*, pages 248–260. PMLR.

Guangzhi Xiong, Qiao Jin, Zhiyong Lu, and Aidong Zhang. 2024. Benchmarking retrieval-augmented generation for medicine. *Preprint*, arXiv:2402.13178.